

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
data=pd.read_csv("abalone.csv")
```

```
data.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
data.shape
```

```
(4177, 9)
```


Univariate analysis

```
sns.boxplot(data.Length)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass t
FutureWarning
```

```
sns.countplot(x = 'Sex', data = data, palette = 'Set3')
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f26821807d0>
```



Bivariate analysis

```
sns.barplot(x=data.Height,y=data.Diameter)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f2682167990>
```




```
numerical_features = data.select_dtypes(include = [np.number]).columns
categorical_features = data.select_dtypes(include = [np.object]).columns
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DeprecationWarning: `np
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/rele
```

```
plt.figure(figsize = (20,7))
sns.heatmap(data[numerical_features].corr(),annot = True)
```

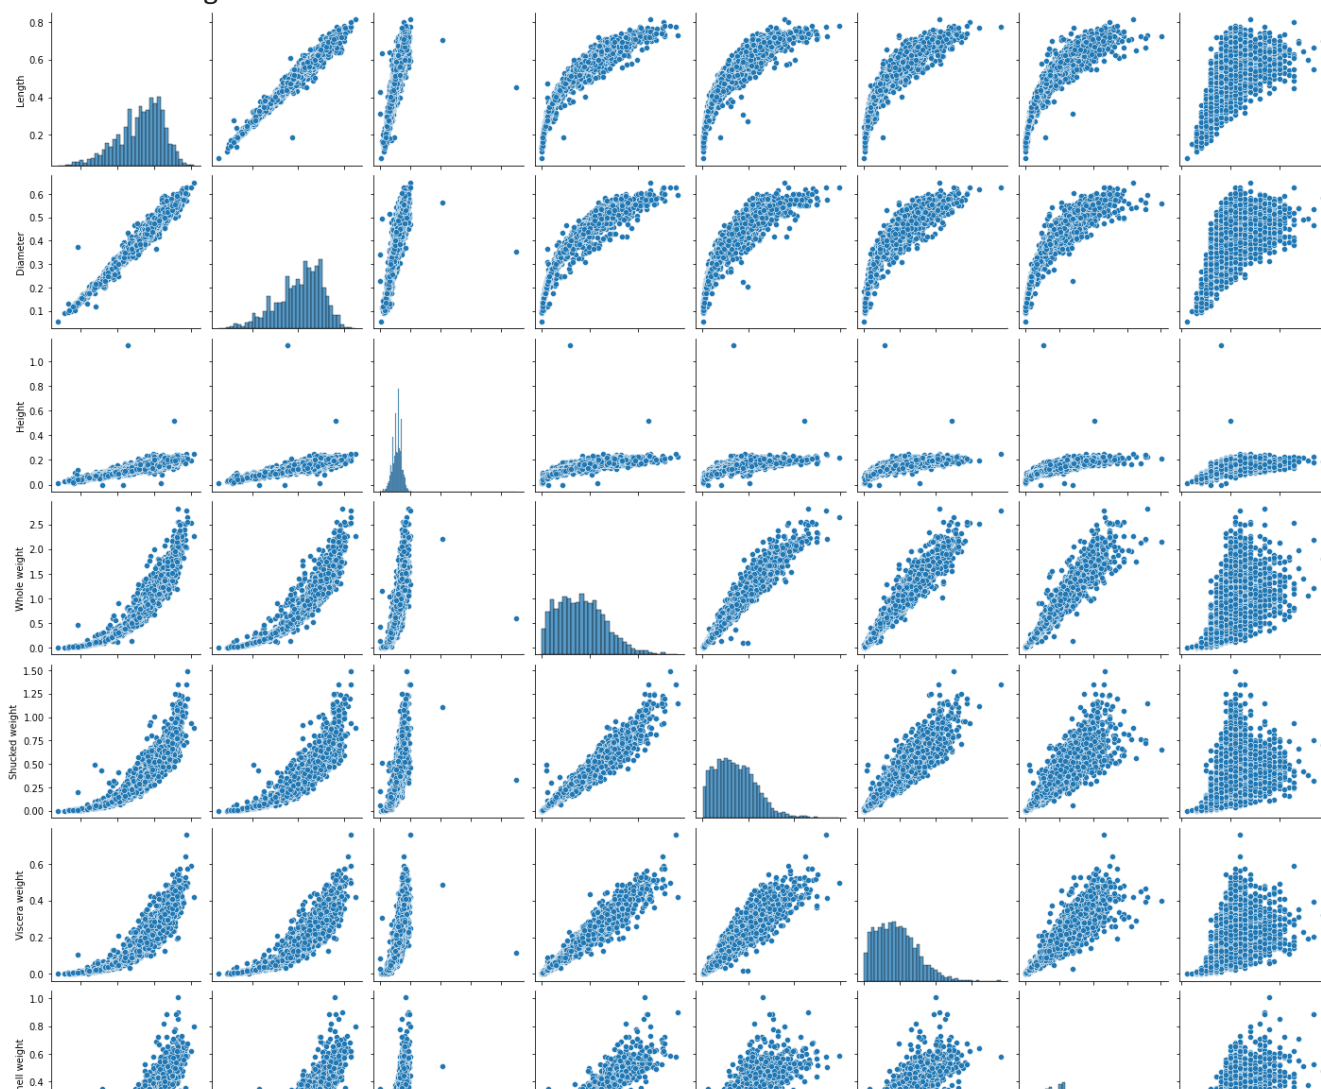
<matplotlib.axes._subplots.AxesSubplot at 0x7f268401db50>



Multivariate analysis

```
sns.pairplot(data)
```

<seaborn.axisgrid.PairGrid at 0x7f26843078d0>



▼ Performing descriptive statistics on the dataset.



```
data['Height'].describe()
```

```
count    4177.000000
mean      0.139516
std       0.041827
min       0.000000
25%       0.115000
50%       0.140000
75%       0.165000
max       1.130000
Name: Height, dtype: float64
```

```
data['Height'].mean()
```

```
0.13951639932966242
```

```
data.max()
```

```
Sex          M
Length      0.815
Diameter    0.65
Height      1.13
Whole weight 2.8255
Shucked weight 1.488
Viscera weight 0.76
Shell weight 1.005
Rings       29
dtype: object
```

```
data['Sex'].value_counts()
```

```
M    1528
I    1342
F    1307
Name: Sex, dtype: int64
```

```
data[data.Height == 0]
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
1257	I	0.430	0.34	0.0	0.428	0.2065	0.0860	0.1150	8
3996	I	0.315	0.23	0.0	0.134	0.0575	0.0285	0.3505	6

```
data['Shucked weight'].kurtosis()
```

```
0.5951236783694207
```

```
data['Diameter'].median()
```

```
0.425
```

```
data['Shucked weight'].skew()
```

```
0.7190979217612694
```

Missing Values

```
data.isna().any()
```

```
Sex          False
Length       False
```

```

Diameter      False
Height         False
Whole weight   False
Shucked weight False
Viscera weight False
Shell weight   False
Rings          False
dtype: bool

```

```

missing_values = data.isnull().sum().sort_values(ascending = False)
percentage_missing_values = (missing_values/len(data))*100
pd.concat([missing_values, percentage_missing_values], axis = 1, keys= ['Missing values', '%

```

	Missing values	% Missing	
Sex	0	0.0	
Length	0	0.0	
Diameter	0	0.0	
Height	0	0.0	
Whole weight	0	0.0	
Shucked weight	0	0.0	
Viscera weight	0	0.0	
Shell weight	0	0.0	
Rings	0	0.0	

Find the outliers

```

q1=data.Rings.quantile(0.25)
q2=data.Rings.quantile(0.75)
iqr=q2-q1
print(iqr)

```


3.0

```

data = pd.get_dummies(data)
dummy_data = data
data.boxplot( rot = 90, figsize=(20,5))

```


<matplotlib.axes._subplots.AxesSubplot at 0x7f267fe90810>



```
data['age'] = data['Rings']
data = data.drop('Rings', axis = 1)
```

```
data.drop(data[(data['Viscera weight'] > 0.5) & (data['age'] < 20)].index, inplace=True)
data.drop(data[(data['Viscera weight'] < 0.5) & (data['age'] > 25)].index, inplace=True)
```

```
var = 'Shell weight'
plt.scatter(x = data[var], y = data['age'])
plt.grid(True)
```



Check for Categorical columns and perform encoding.

```
numerical_features = data.select_dtypes(include = [np.number]).columns
categorical_features = data.select_dtypes(include = [np.object]).columns
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DeprecationWarning: `np` deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/rele>

◀ ▶

```
abalone_numeric = data[['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight', 'Visc
```

```
abalone_numeric.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age	Sex_F	Sex_I	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15	0	0	1
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7	0	0	1
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9	1	0	0
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10	0	0	1

```
x = data.iloc[:, 0:1].values
```

```
y = data.iloc[:, 1]
```

```
y
```

```
0      0.365
```

```
1      0.265
```

```
2      0.420
```

```
3      0.365
```

```
4      0.255
```

```
...
```

```
4172    0.450
```

```
4173    0.440
```

```
4174    0.475
```

```
4175    0.485
```

```
4176    0.555
```

```
Name: Diameter, Length: 4150, dtype: float64
```

Scale the independent variables

```
print ("ORIGINAL VALUES: \n", x,y)
```

```
ORIGINAL VALUES:
```

```
[[0.455]
```

```
[0.35 ]
```

```
[0.53 ]
```

```
...
```

```
[0.6  ]
```

```
[0.625]
```

```
[0.71 ]] 0      0.365
```

```
1      0.265
```

```
2      0.420
```

```
3      0.365
```

```
4      0.255
```

```
...
```

```
4172    0.450
```

```
4173    0.440
```

```
4174    0.475
```



```

4175    0.485
4176    0.555
Name: Diameter, Length: 4150, dtype: float64

```

```

from sklearn import preprocessing
min_max_scaler = preprocessing.MinMaxScaler(feature_range =(0, 1))
new_y= min_max_scaler.fit_transform(x,y)
print ("VALUES AFTER MIN MAX SCALING: \n", new_y)

```

```

VALUES AFTER MIN MAX SCALING:
[[0.51351351]
 [0.37162162]
 [0.61486486]
 ...
 [0.70945946]
 [0.74324324]
 [0.85810811]]

```

Split the data into training and testing

```

X = data.drop('age', axis = 1)
y = data['age']

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_selection import SelectKBest
standardScale = StandardScaler()
standardScale.fit_transform(X)

selectkBest = SelectKBest()
X_new = selectkBest.fit_transform(X, y)

X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size = 0.25)
X_train

```

```

array([[0.73 , 0.55 , 0.205, ..., 1.    , 0.    , 0.    ],
       [0.395, 0.3   , 0.09 , ..., 0.    , 1.    , 0.    ],
       [0.625, 0.495, 0.175, ..., 0.    , 0.    , 1.    ],
       ...,
       [0.645, 0.51 , 0.18 , ..., 1.    , 0.    , 0.    ],
       [0.71 , 0.56 , 0.18 , ..., 1.    , 0.    , 0.    ],
       [0.46 , 0.36 , 0.135, ..., 0.    , 0.    , 1.    ]])

```

Build the Model

```

from sklearn import linear_model as lm
from sklearn.linear_model import LinearRegression
model=lm.LinearRegression()
results=model.fit(X_train,y_train)

```

```
accuracy = model.score(X_train, y_train)
print('Accuracy of the model:', accuracy)
```

Accuracy of the model: 0.5305010084253585



Train the Model

Train the Model

```
lm = LinearRegression()
lm.fit(X_train, y_train)
y_train_pred = lm.predict(X_train)
y_train_pred
```

array([19.9375, 6.9375, 10.625 , ..., 10.9375, 12.125 , 11.6875])

X_train

```
array([[0.73 , 0.55 , 0.205, ..., 1.    , 0.    , 0.    ],
       [0.395, 0.3   , 0.09 , ..., 0.    , 1.    , 0.    ],
       [0.625, 0.495, 0.175, ..., 0.    , 0.    , 1.    ],
       ...,
       [0.645, 0.51 , 0.18 , ..., 1.    , 0.    , 0.    ],
       [0.71 , 0.56 , 0.18 , ..., 1.    , 0.    , 0.    ],
       [0.46 , 0.36 , 0.135, ..., 0.    , 0.    , 1.    ]])
```

y_train

```
3081    14
3602     5
1805     9
3898    10
2555     6
..
3913    11
594     12
1955    12
1196    11
2484    14
Name: age, Length: 3112, dtype: int64
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error
s = mean_squared_error(y_train, y_train_pred)
print('Mean Squared error of training set :%2f'%s)
```

Mean Squared error of training set :4.775387

Test the Model

```
y_train_pred = lm.predict(X_train)
y_test_pred = lm.predict(X_test)
y_test_pred

array([10.875 , 10.375 ,  9.8125, ..., 12.6875, 10.375 ,  8.0625])

p = mean_squared_error(y_test, y_test_pred)
print('Mean Squared error of testing set :%2f'%p)

Mean Squared error of testing set :4.556061
```

Measure the performance using Metrics.

```
from sklearn.metrics import r2_score
s = r2_score(y_train, y_train_pred)
print('R2 Score of training set:%.2f'%s)

R2 Score of training set:0.53

p = r2_score(y_test, y_test_pred)
print('R2 Score of testing set:%.2f'%p)

R2 Score of testing set:0.55
```

✓ 0s completed at 5:36 PM

